I love AI. Why doesn't everyone?

Anti-AI sentiment might or might not be rational, but it certainly relies on a lot of bad arguments.

NOAH SMITH
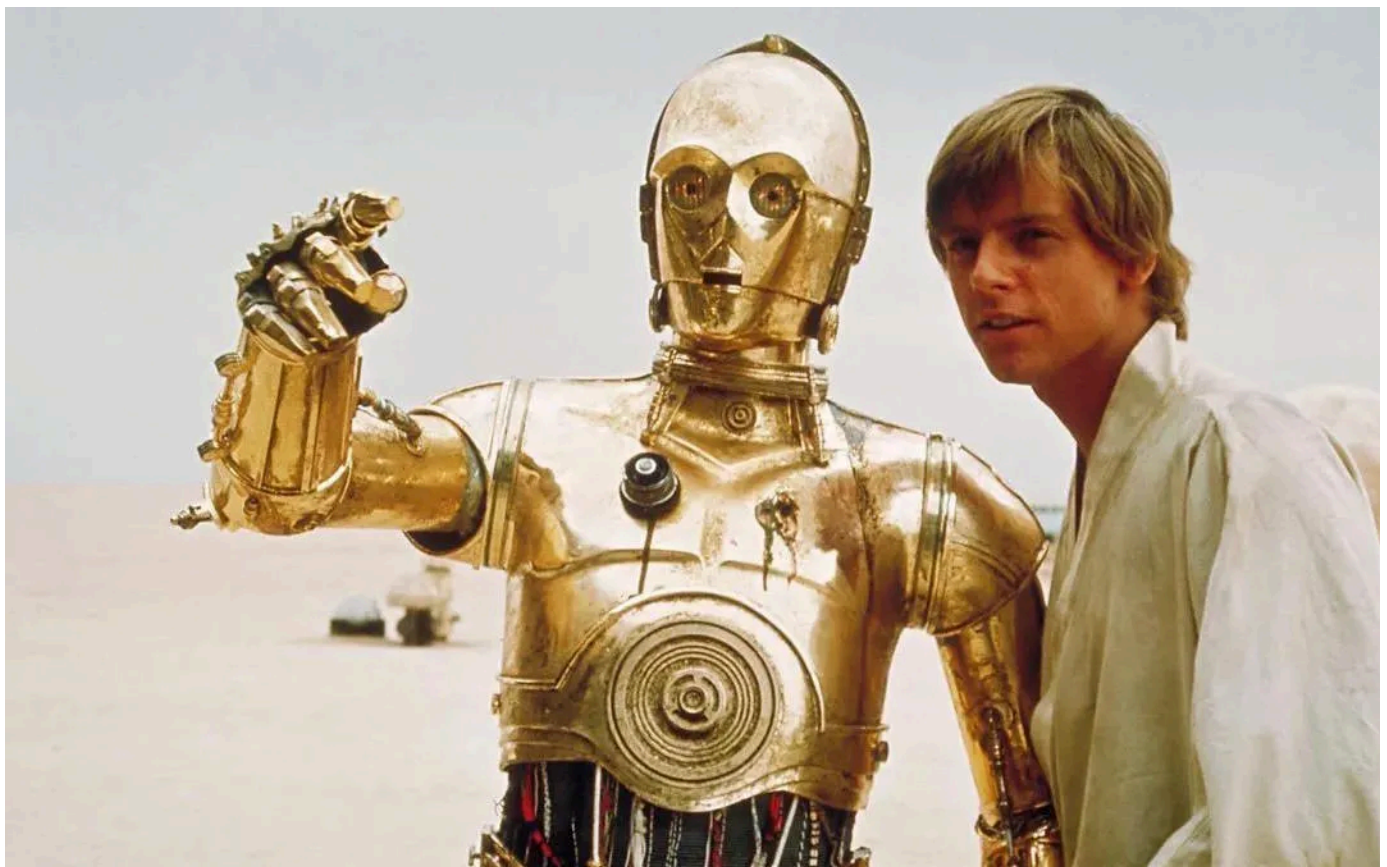DEC 01, 2025

♡ 434   💬 324   🔁 81   Sha

New technologies almost always create lots of problems and challenges for our soci The invention of farming caused local overpopulation. Industrial technology caused pollution. Nuclear technology enabled superweapons capable of destroying civilization. New media technologies arguably cause social unrest and turmoil whenever they're introduced.

And yet how many of these technologies can you honestly say you wish were never invented? Some people romanticize hunter-gatherers and medieval peasants, but I don't see many of them rushing to go live those lifestyles. I myself buy into the argument that smartphone-enabled social media is largely responsible for a variety modern social ills, but I've always maintained that eventually, our social institution will evolve in ways that minimize the harms and enhance the benefits. In general, when we look at the past, we understand that technology has almost always made things better for humanity, especially over the long haul.
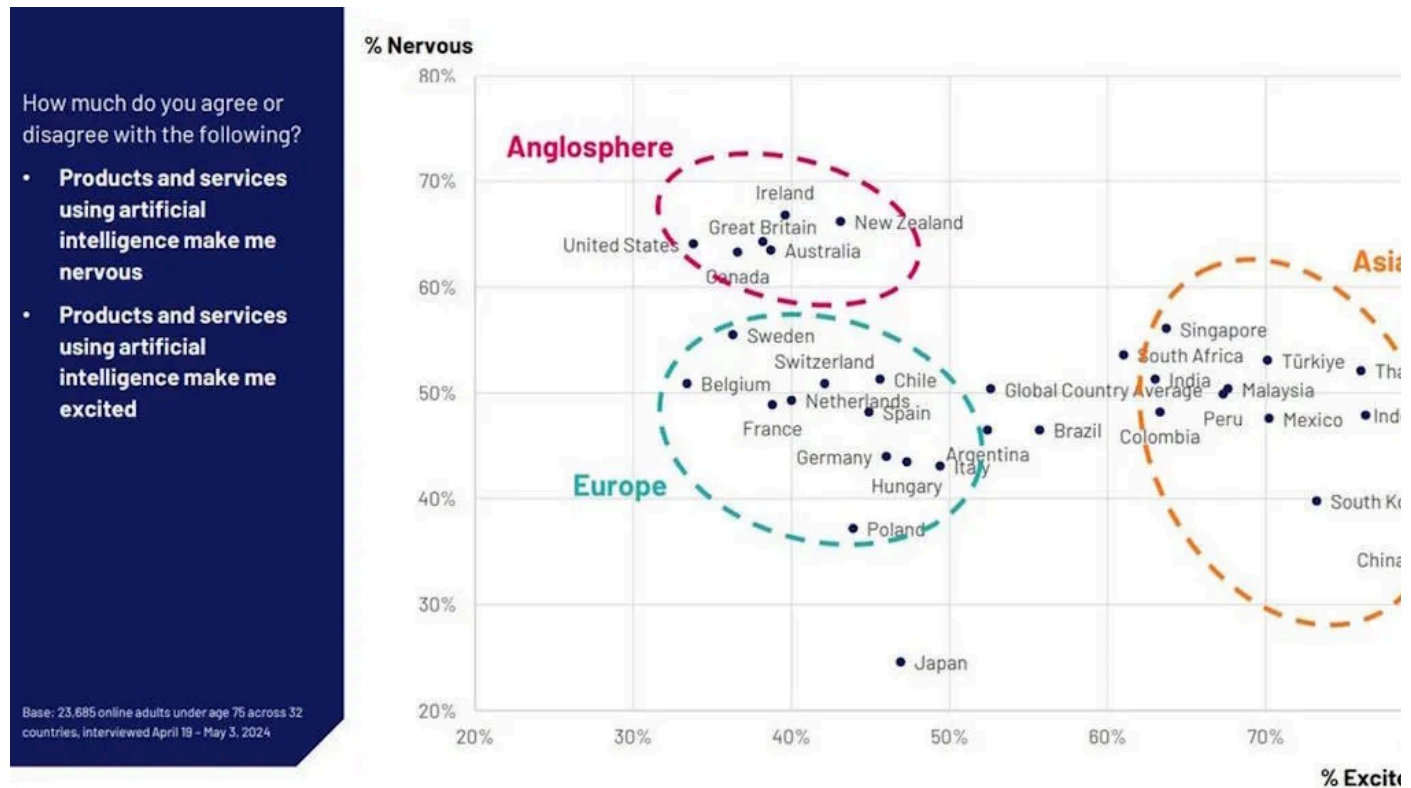
But when we think about the technologies *now being invented*, we often forget this lesson — or at least, many of us do. In the U.S., there have recently been movement against mRNA vaccines, electric cars, self-driving cars, smartphones, social media, nuclear power, and solar and wind power, with varying degrees of success.

The difference between our views of old and new technologies isn't necessarily irrational. Old technologies present less risk — we basically know what effect they' have on society as a whole, and on our own personal economic opportunities. New technologies are disruptive in ways we can't predict, and it makes sense to be worri about that risk that we might personally end up on the losing end of the upcoming social and economic changes.

But that still doesn't explain changes in our attitudes toward technology *over time*. Americans largely embraced the internet, the computer, the TV, air travel, the automobile, and industrial automation. And risk doesn't explain all of the differenc in attitudes among *countries*.

In the U.S., few technologies have been on the receiving end of as much popular fea and hatred as generative AI. Although policymakers have remained staunchly in fav of the technology — probably because it's supporting the stock market and the economy — regular Americans of both parties tend to say they're more concerned t excited, with an especially rapid increase in negative sentiment among progressives

There is plenty of trepidation about AI around the world, but America stands out. A 2024 Ipsos poll found that no country surveyed was both more nervous and less exc about AI than the United States:
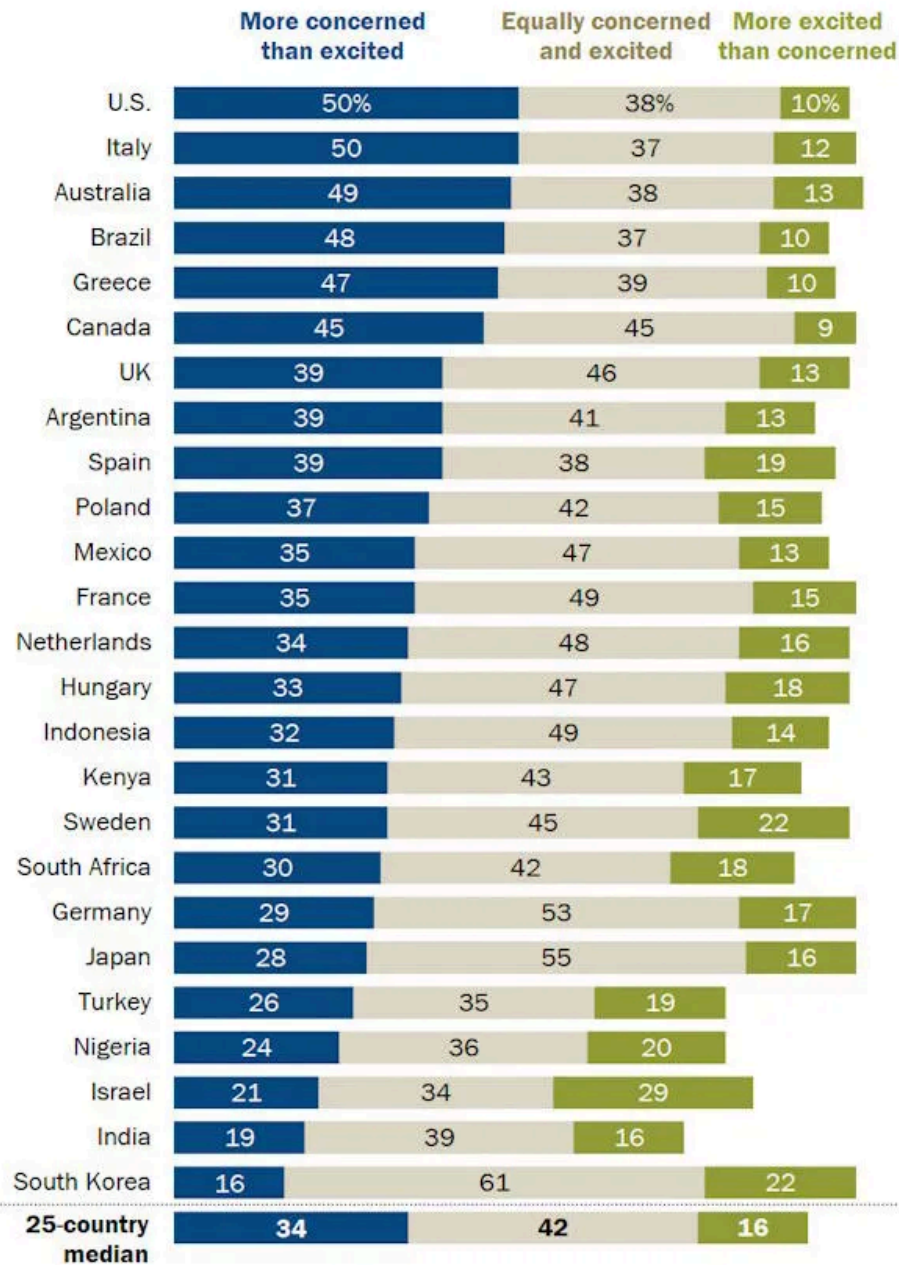

Source: Ipsos

America's fear of AI stands in stark contrast to countries in Asia, from developing countries like India and Indonesia to rich countries like South Korea and Singapore Even Europe, traditionally not thought of as a place that embraces the new, is significantly less terrified than the U.S. Other polls find similar results:

How do people around the world feel about the rise of AI in daily life?

% who say the increased use of artificial intelligence in daily life makes them feel ...

	More concerned than excited	Equally concerned and excited	More excited than concerned
U.S.	50%	38%	10%
Italy	50	37	12
Australia	49	38	13
Brazil	48	37	10
Greece	47	39	10
Canada	45	45	9
UK	39	46	13
Argentina	39	41	13
Spain	39	38	19
Poland	37	42	15
Mexico	35	47	13
France	35	49	15
Netherlands	34	48	16
Hungary	33	47	18
Indonesia	32	49	14
Kenya	31	43	17
Sweden	31	45	22
South Africa	30	42	18
Germany	29	53	17
Japan	28	55	16
Turkey	26	35	19
Nigeria	24	36	20
Israel	21	34	29
India	19	39	16
South Korea	16	61	22
25-country median	**34**	**42**	**16**

Note: Those who did not answer are not shown.
Source: Spring 2025 Global Attitudes Survey.
"How People Around the World View AI"

PEW RESEARCH CENTER

Source: [Pew](#)

If Koreans, Indians, Israelis, and Chinese people aren't terrified of AI, why should Americans be so scared — especially when we usually embraced previous technolog wholeheartedly? Do we know something they don't? Or are we just biased by some combination of political unrest, social division, wealthy entitlement, and disconnection from physical industry?

It's especially dismaying because I've spent most of my life dreaming of having something like modern AI. And now that it's here, I (mostly) love it.

I always wanted a little robot friend, and now I have one

Media has prepared me all my life for AI. Some of the portrayals were negative, of course — Skynet, the computer in the Terminator series, tries to wipe out humanity and HAL 9000 in *2001: A Space Odyssey* tries to kill its user. But most of the AIs depicted in sci-fi were friendly — if often imperfect — robots and computers.

C-3PO and R2-D2 from Star Wars are Luke's loyal companions, and save the Rebell on numerous occasions — even if C-3PO is often wrong about things. The ship's computer in Star Trek is a helpful, reassuring presence, even if it occasionally mess up its holographic creations. [1] Commander Data from *Star Trek: The Next Generation* a heroic figure, probably based on a character from Isaac Asimov's Robot series — is just one of hundreds of sympathetic portrayals of androids. Friendly little rolling robots like Wall-E and Johnny 5 from *Short Circuit* are practically a stock character, helpful sentient computers are important protagonists in *The Moon is a Harsh Mistre* the Culture novels, the TV show *Person of Interest*, and so on. The novel *The Diamon Age* features an AI tutor that helps kids out of poverty, while the *Murderbot* series is about a security robot who just wants to live in peace.

In these portrayals, intelligent robots and computers are consistently portrayed as helpful assistants, allies, and even friends. Their helpfulness makes sense, since the created to be our tools. But some deep empathetic instinct in our human nature ma

it difficult to objectify something so intelligent-seeming as a simple tool. And so it'
natural for us to portray AIs as *friends*.

Fast forward a few decades, and I actually have that little robot friend I always
dreamed of. It's not exactly like any of the AI portrayals from sci-fi, but it's
recognizably similar. As I go through my daily life, GPT (or Gemini, or Claude) is
always there to help me. If my water filter needs to be replaced, I can ask my robot
friend how to do it. If I forget which sociologist claimed that economic growth crea
the institutional momentum for further growth, [2] I can ask my robot friend who tha
was. If I want to know some iconic Paris selfie spots, it can tell me. If I can't remem
the article I read about China's innovation ecosystem last year, my robot buddy can
find it for me.

It can proofread my blog posts, be my search engine, help me decorate my room,
translate other languages for me, teach me math, explain tax documents, and so on.
This is just the beginning of what AI can do, of course. It's possibly the most genera
purpose technology ever invented, since its basic function is to memorize the entire
corpus of human knowledge and then spit any piece of it back to you on command.
And because it's programmed to do everything with a smile, it's always friendly and
cheerful — just like a little robot friend ought to be.

No, AI doesn't always get everything right. It makes mistakes fairly regularly. But I
never expected engineers to be able to create some kind of infallible god-oracle tha
knows every truth in the Universe. C-3PO gets stuff confidently wrong all the time,
does the computer on Star Trek. For that matter, so does my dad. So does every hun
being I've ever met, and every news website I've ever read, and every social media
account I've ever followed. Just like with every other source of information and
assistance you've ever encountered in your life, AI needs to be cross-checked before
you can believe 100% in whatever it tells you. Infallible omniscience is still beyond t
reach of modern engineering.

Who cares? This is an amazingly useful technology, and I love using it. It has opene[
my informational horizons by almost as much as the internet itself, and made my lit
far more convenient. Even without the expected impacts on productivity, innovatio
and so on, just having this little robot friend would be enough for me to say that AI
improved my life enormously.

This instinctive, automatic reaction to such a magical new tool seems utterly natura
to me. And yet when I say this on social media, people pop out of the woodwork to
denounce AI and ridicule anyone who likes it. Here are just a few examples:

Normally I would just dismiss these outbursts as non-representative. But in this cas
there's pretty robust survey data showing that the American public is overwhelming
negative on AI. These social media malcontents may be unusually vituperative, but
their opinions probably aren't out of the mainstream.

What's going on here? Why doesn't everyone else love having a little robot friend w
can answer lots of their questions and perform lots of their menial tasks?

I guess it makes sense that for a lot of people, the potential negative externalities —
deepfakes, the decline of critical thinking, ubiquitous slop, or the risk that bad acto

will be able to use AI to do major violence — loom large. Other people, like artists
translators, may fear for their careers. I think it's likely that in the long run, our soc
will learn to deal with all those challenges, but as Keynes said, "in the long run we'r
all dead."

And yet the instinctive negativity with which AI is being met by a large segment of
American public feels like an unreasonable reaction to me. Although externalities a
distributional disruptions certainly exist, the specific concerns that many of AI's m
strident critics cite are often nonsensical.

A lot of the anti-AI canon is nonsense

One of the most common talking points you hear about AI is that data centers use a
ton of water, potentially causing water shortages. For example, *Rolling Stone* recentl
put out an article by Sean Patrick Cooper, entitled "The Precedent Is Flint: How
Oregon's Data Center Boom Is Supercharging a Water Crisis". Here's what it claim

> [D]ata centers pose a variety of climate and environmental problems, including t
> impact on the water supply. The volume of water needed to cool the servers in da
> centers — most of which need to be kept at 70 to 80 degrees to run effectively —
> become a nationwide water resource issue particularly in areas facing water scar
> across the West. This year, a Bloomberg News analysis found that roughly "two-
> thirds of new data centers built or in development since 2022 are in places alread
> gripped by high levels of water stress." Droughts have plagued Morrow County,
> occurring annually since 2020. But even areas with ample water reserves are
> vulnerable to the outsized demand from data centers. Earlier this year, the
> International Energy Agency reported that data centers could consume 1,200 bill
> liters by 2030 worldwide, nearly double the 560 billion liters of water they use
> currently.

The idea that AI data centers are water-guzzlers has become standard canon in mar
areas of the internet — especially among progressives. And yet it's just not a real is

Andy Masley finally got fed up and dug into the data, writing an epic blog post that debunked every single version of the "AI uses lots of water" narrative:

> 🐢 **The Weird Turn Pro**
>
> ## The AI water issue is fake
>
> AI data centers use water. Like any other industry that uses water, they require careful planning. If an electric car factory opens near you, that factory may use just as much water as a data center. The factory also requires careful planning. But the idea that either the factory or AI is using an inordinate amount of water that merits any kind of boyco...
>
> [Read more]
>
> 2 months ago · 160 likes · 30 comments · Andy Masley

Masley notes that A) almost all of the water AI uses is actually used by power plants generating electricity to power AI, and B) most of the water that gets "used" by AI is actually just run through the system and returned to the original source, instead of being used up. He writes:

> So of the ways AI uses water...The vast majority (maybe 90%) is withdrawn, freshwater (not potable) that is indirectly (offsite) used non-consumptively in power plants (it's returned to the source unaffected)...Less (maybe 7%) is withdrawn freshwater (not potable) that is consumed (evaporated) indirectly (offsite) in the power plants to generate the electricity AI uses...And less (maybe 3%) is withdrawn freshwater that's then treated to become potable, used directly (onsite) in physical data centers themselves, and consumed after (not returned to the source, evaporated).

He then goes on to pull a bunch of numbers that show that AI's actual water consumption isn't a problem (yet):
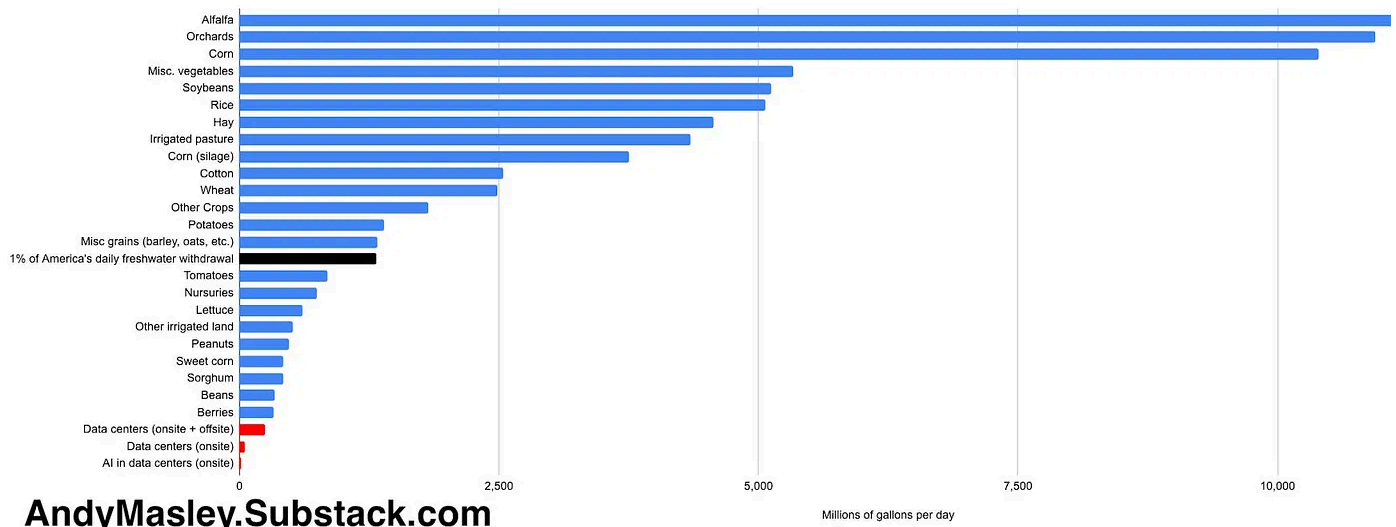
> All U.S. data centers (which mostly support the internet, not AI) used 200–250 million gallons of freshwater daily in 2023. The U.S. consumes approximately 13: billion gallons of freshwater daily…So data centers in the U.S. consumed approximately 0.2% of the nation's freshwater in 2023…However, the water that v actually used onsite in data centers was only 50 million gallons per day…Only 0.0 of America's freshwater in 2023 was consumed inside data centers themselves. T is 3% of the water consumed by the American golf industry.
>
> AI uses approximately 20% of the electricity in data centers…Water use roughly correlates with electricity…So AI consumes approximately 0.04% of America's freshwater if you include onsite and offsite use, and only 0.008% if you include ju the water in data centers. So AI…is using 0.008% of America's total freshwater…
>
> So the water all American data centers will consume onsite in 2030 is equivalent
>
> - 8% of the water currently consumed by the U.S. golf industry.
> - 8% of the water consumed by U.S. steel production.
> - The water usage of 260 square miles of irrigated corn farms, equivalent to 1% America's total irrigated corn.

And he includes a chart for comparison to other water uses:

Source: <u>Andy Masley</u>

He goes on to debunk other aspects of the "AI water use" story, showing that data centers haven't hurt water availability in local areas, and that data centers aren't a major source of pollution. After that, he goes through some articles claiming a link between AI and water use, and finds that either their data is obviously wrong, or th don't even cite data.

Masley is especially annoyed with a book by Karen Hao called *Empire of AI*, which made huge math errors about AI water use:

> Within 20 pages, Hao manages to:
>
> Claim that a data center is using 1000x as much water as a city of 88,000 people, where it's actually using about 0.22x as much water as the city, and only 3% of th municipal water system the city relies on. She's off by a factor of 4500.
>
> Imply that AI data centers will consume 1.7 trillion gallons of drinkable water by 2027, while the study she's pulling from says that only 3% of that will be drinkabl water, and 90% will not be consumed, and instead returned to the source unaffec
>
> Paint a picture of AI data centers harming water access in America, where they don't seem to have caused any harm at all.

Hao later admitted some serious data errors in her book.

Anyway, Masley's whole post is very long and involved, and it links to a bunch of ot
long and involved posts on the various subtopics. If you're interested in the "AI wat
use" issue at all, this is a must-read. I'm not sure I've ever seen so thorough a
debunking of a popular belief. Some AI critics like Timnit Gebru attacked Masley,
were not able to muster any substantive rebuttal; instead, they simply suggested tha
Masley speak to activists:

There are two things that are especially frustrating about the "AI and water use"
argument. First, there's the powerlessness of facts in the face of viral misinformatic
If a bogus claim gets repeated by enough people, it becomes its own sort of echo
chamber, where observers believe the false claim because they think they've heard i
from multiple sources. Politically motivated reasoning only adds to that effect — a
of Americans fear AI and want to find a reason why it's bad, so they latch onto the
"water use" story without checking to see if it's real.

The second frustrating thing is that *there's a much better argument sitting right there for the taking.* AI doesn't use up much water, and probably won't for the foreseeable future. But it does use a hell of a lot of *electricity*, and this could pose a problem if the tech keeps scaling up. That electricity use can strain local grids and raise carbon emissions. It's a real challenge! And yet AI critics tend to ignore this grounded and reasonable worry in favor of the "water use" myth.

Nor is this the only such example. Aaron Regunberg, writing in *The New Republic*, trots out a vast litany of accusations against AI, most of which don't hold up under scrutiny. For example, he claims that an AI crash would wipe out normal Americans while preserving the wealth of AI's creators:

> First, there is the likelihood that the AI industry is building up a bubble that, when it bursts, will take down the global economy…When this thing pops, it won't be the filthy rich scammers behind the bubble who will lose out. It will be regular people. One former International Monetary Fund chief economist estimates that a crash could wipe out $20 trillion in wealth held by American households.

But that's absurd. Gita Gopinath's analysis, which Regunberg cites, is all about *stock* wealth. Most of the wealth of the people who create AI is in their own company's stocks, so they would *absolutely* get wiped out in a crash. Regular Americans would hurt a bit, but most of their wealth is in their houses. Almost all American stocks are owned by the rich. In fact, there was a tech stock crash in 2022, and wealth inequality went down as a result.

Anyway, Regunberg continues:

> Beyond the profoundly compelling politics of a market crash and bailout, there's the equally potent issue of AI-driven job losses. These hits have already begun. Last summer, IBM replaced hundreds of employees with AI chatbots, and UPS, JPMorgan Chase, and Wendy's have all begun following suit. The CEO of Anthropic has warned that AI could lead to the loss of half of all entry-level white

> collar jobs; McKinsey estimated that AI could automate 60 to 70 percent of
> employees' work activities; and a recent report from Senator Bernie Sanders and
> Health, Education, Labor and Pensions Committee found that AI could replace &
> percent of fast-food and counter workers, 64 percent of accountants, and 47 perc
> of truck drivers. The anger that such dislocations will generate against "clankers
> yes, anti-AI frustration has already inspired the explosion of a Gen Z meme-slur
> hard to overstate, potentially exceeding the resentment against the North Ameri
> Free Trade Agreement that Trump rode to the White House in 2016.

The studies Regunberg cites, predicting how many jobs will be "replaced", actually
no such thing, as I explained in this post back in 2023:

Nobody knows how many jobs will "be automated"
NOAH SMITH • APRIL 10, 2023
Read full story →

They are merely engineers' guesses about which jobs will be *affected* by AI.

In fact, one recent study found that industries that are predicted to use AI more are
seeing no slowdown in wages, and have experienced robust employment growth for
workers in their 30s, 40s, and 50s:

AI and jobs, again
NOAH SMITH • AUG 30
Read full story →

The study did find a slowdown in hiring for younger workers. But other studies find
effect of AI on jobs at all so far. Obviously, many workers are afraid of losing their j
to AI, but those losses don't seem to have materialized yet, so it's a bit silly when A
critics like Regunberg present job losses as established fact.

How can folks like Regunberg, Gebru, Hao, and Cooper get away with such incredil
sloppiness? Unfortunately, the answer is "motivated reasoning". So many American
feel so negatively about AI that if some writer or intellectual comes up and makes

some claim that AI is a water-guzzler or is causing mass unemployment, some peor will believe it. Fear and anger have a way of finding their own justifications.

So I don't think there's much I can do about the unreasonable negative emotions be directed at AI in the United States. I've simply got to reconcile myself to the fact th my delight in this new technology is a rare minority position. I do think there are probably some things that AI companies can do to improve their image with the American public. But for now, all I can do is lament that a country that once boldly embraced the future now quails at it.

| Type your email... | Subscribe |

1   In fact, I'd argue that the Star Trek computer is the sci-fi portrayal of AI that most accurately predicts modern generative AI. Whole plots of *Star Trek: The Next Generation* revolve around holodeck prompts that generate hallucinations.

2   A.O. Hirshman

434 Likes · 81 Restacks

← Previous   Next →

Discussion about this post

Comments   Restacks

Write a comment...

Matthew 🕸 5d

The droids in star wars are literal slaves. Like that is their purpose. We meet R2D2 and C3P0 they are sold by the jawas to Luke's uncle.

I am being a bit facetious here, but not pointing that out in the post seems like an oversight.

Also, this post complains the most about the false belief in water loss from AI. There are many better things to complain about vis a vis AI, but the post hides behind the water thing to avoid having to go into depth about the others.

♡ LIKE (109)   💬 REPLY   ⬆

27 replies by Noah Smith and others

What-username-999 🕸 5d

I don't hate AI. I do hate that companies, mine included, are giving mandates to incorporate it work without actually thinking it through. I hate the tech bro hype surrounding it. I hate the ca class will try and use it to immiserate everyone else to increase their wealth.

If it crashed and took all of their wealth and knocked tech bros down several pegs, I'd crack a to be honest.

♡ LIKE (86)   💬 REPLY   ⬆

57 replies by Noah Smith and others

322 more comments…